




The Secure Evidence and Credential Trust (SECT) Platform

.. ——— The Architecture of Defensible Decisions ——— ..

SUBMISSION OVERLOAD



Review teams process 40 to 500 submissions per role or intake period, creating manual bottlenecks.



INTEGRITY GAPS

EXPERIENCE

SKILLS

DEGREE

CERTIFICATIONS

YEARS

TITLE



Keyword filters miss context, failing to catch invalid, inflated, or incomplete submissions before human review.

THE RESULT: Increased legal risk, non-explainable filtering, and critical compliance exposure. High-stakes credentialing requires a system built for audit-grade defensibility, not basic applicant tracking.

Introducing the Trust Shell



Human-in-the-Loop
Equips authorized personnel with high-speed clarity, reserving all high-impact decision-making for human evaluators.



Not a Lie Detector


SECT mathematically measures the alignment between a candidate's claims, their provided evidence, and institutional policy.



Cryptographic Shield

Every automated evaluation and human intervention is secured with an immutable, traceable cryptographic seal.


The 5-Layer Operational Architecture




A modular architecture separating raw evidence from evaluation logic ensures repeatable, scrubbable, and legally defensible results.

Layer 1: The Evidence Layer (From Chaos to Fact)


Normalization is the prerequisite for fairness. By converting unstructured PDFs into a structured inventory, every candidate is measured against the exact same baseline, removing the noise of formatting.



Layer 2: The Policy Layer (The Eligibility Gate Engine)



Layer 3: Multi-Pass Decision Scoring




Decision Synthesis: The Output Semantics Matrix

Output Class	Definition & Reason Codes	System Action
 Advance	 Meets all policy requirements with high confidence.	 Routes to Shortlist.
Hold	 Incomplete submission or minor chronology inconsistency.	 Requires manual clarification.
Escalate	 High-risk discrepancies or claim-evidence mismatch codes.	 Routes to Security Lead Escalation Queue.
Reject	 Fails mandatory criteria (Hard disqualification code).	 Blocks approval pending human ratification.

 **Key Takeaway:** Every automated output must be accompanied by a specific reason code and policy citation. No black box outcomes.


Layer 4: Workflow & Stakeholder Routing



Accountability Framework: Overrides & The Two-Person Rule


The Two-Person Rule

Any automated outcome designated as Disqualified or Security-Escalation acts as a hard voter requirement, blocking state transitions until a second authorized human signature is logged.



Full Override Accountability

Human reviewers have total authority to override machine suggestions, but their interventions are strictly governed.



Layer 5: The Defensibility Shield (Merkle Integrity)


The Verity Engine:

Generates a unique cryptographic hash for every artifact in a case dossier.


By computing a single Merkle root, it creates an immutable linkage between a decision and its underlying evidence.

Self-Authenticating Evidence:

If a document is altered retrospectively, the Merkle root breaks, ensuring absolute evidentiary weight during legal disputes.



Ethical Guardrails & Non-Negotiable Principles



Pillar 1: No Autonomous Adverse Decisions.

Algorithms never reject a candidate autonomously; humans must ratify all disqualifications.

Pillar 2: Universal Explainability.

Every score and class must output a specific, traceable reason code. No opaque filtering.

Pillar 3: Least Privilege Access.

Role-based data access mitigates the blast radius of potential breaches. Protected attributes are excluded from fit-scoring.


Pillar 4: Policy Version Locking.

Cases are evaluated against the specific rules active at submission time, preventing retrospective compliance drift.


Pillar 5: Override Traceability.

Human deviations are logged with the same cryptographic rigor as machine logic.

Adaptability: Sector-Specific Configuration Packs



Business Impact & Operational KPIs



The 30-Day Proof of Value

Objective: Demonstrate faster, more defensible shortlist decisions on a single hiring stream of 100 to 300 submissions—without relinquishing human control.

